

Logstash Integration


ZABBIX

+



Origins

- ▶ Jordan Sissel
- ▶ Started in 2009
- ▶ Open Source (Apache License)
- ▶ Jordan joined **Elastic** in August 2013
- ▶ Still Open Source
- ▶ Will always be Open Source



What is it?

- ▶ A tool for receiving, processing and outputting logs, and other data streams.
- ▶ Pipeline
 - ▶ Input
 - ▶ Filter
 - ▶ Output



Inputs

- couchdb_changes
- drupal_dblog
- elasticsearch
- exec
- eventlog
- file
- ganglia
- gelf
- generator
- graphite
- github
- heartbeat
- heroku
- http
- http_poller
- irc
- imap
- jdbc
- jmx
- kafka
- log4j
- lumberjack
- meetup
- pipe
- puppet_facter
- relp
- rss
- rackspace
- rabbitmq
- redis
- snmptrap
- stdin
- sqlite
- s3
- sqs
- stomp
- syslog
- tcp
- twitter
- unix
- udp
- varnishlog
- wmi
- websocket
- xmpp
- zenoss
- zeromq

Filters

- aggregate
- alter
- anonymize
- collate
- csv
- cidr
- clone
- cipher
- checksum
- date
- dns
- drop
- elasticsearch
- extractnumbers
- environment
- elapsed
- fingerprint
- geoip
- grok
- i18n
- json
- json_encode
- kv
- mutate
- metrics
- multiline
- metaevent
- prune
- punct
- ruby
- range
- syslog_pri
- sleep
- split
- throttle
- translate
- uuid
- urldecode
- useragent
- xml
- zeromq

Outputs

- boundary
- circonus
- csv
- cloudwatch
- datadog
- datadog_metrics
- email
- elasticsearch
- exec
- file
- google_bigquery
- google_cloud_storage
- ganglia
- gelf
- graphtastic
- graphite
- hipchat
- http
- irc
- influxdb
- juggernaut
- jira
- kafka
- lumberjack
- librato
- loggly
- mongodb
- metriccatcher
- nagios
- null
- nagios_nsc
- opentsdb
- pagerduty
- pipe
- riemann
- redmine
- rackspace
- rabbitmq
- redis
- riak
- s3
- sqs
- stomp
- statsd
- solr_http
- sns
- syslog
- stdout
- tcp
- udp
- webhdfs
- websocket
- xmpp
- **zabbix**
- zeromq

Configuration

```
input {  
    plugin_name { settings... }  
}  
filter {  
    plugin_name { settings... }  
}  
output {  
    plugin_name { settings... }  
}
```

Inputs

file

Read events from a file in real-time,
like `tail`

file

```
file {  
  path => "/path/to/logfile"  
}
```

tcp

Read from TCP socket

tcp

```
tcp {  
  host => "ip or hostname"  
  port => 12345  
}
```

irc

Capture all or part of the discussion in one or more IRC channels.

irc

```
irc {  
    channels => [ "#zabbix" ]  
    host => "irc.freenode.org"  
    nick => "my_nickname"  
    port => 6667  
}
```

Inputs

- couchdb_changes
- drupal_dblog
- elasticsearch
- exec
- eventlog
- file
- ganglia
- gelf
- generator
- graphite
- github
- heartbeat
- heroku
- http
- http_poller
- irc
- imap
- jdbc
- jmx
- kafka
- log4j
- lumberjack
- meetup
- pipe
- puppet_facter
- relp
- rss
- rackspace
- rabbitmq
- redis
- snmptrap
- stdin
- sqlite
- s3
- sqs
- stomp
- syslog
- tcp
- twitter
- unix
- udp
- varnishlog
- wmi
- websocket
- xmpp
- zenoss
- zeromq

Filters

grok

Parse arbitrary text and structure it.

grok

- ▶ Parse unstructured log data into something structured.
- ▶ Perfect for syslog, webserver, & db logs, and in general, any log format that is generally written for humans.
- ▶ Ships with 120+ patterns. You can add your own trivially.
- ▶ For help building patterns to match your logs:
 - ▶ <http://grokconstructor.appspot.com/>
 - ▶ <http://grokdebug.herokuapp.com>

grok

55.3.244.1 GET /index.html 15824 0.043

```
filter {  
  grok {  
    match => { "message" => "%{IP:client} %{WORD:method}  
%{URIPATHPARAM:request} %{NUMBER:bytes} %{NUMBER:duration}" }  
  }  
}
```

grok

- ▶ client: 55.3.244.1
- ▶ method: GET
- ▶ request: /index.html
- ▶ bytes: 15824
- ▶ duration: 0.043

grok

Oniguruma

- ▶ `(?<field_name>the pattern here)`
- ▶ `(?<queue_id>[0-9A-F]{10,11})`

Custom patterns_dir

- ▶ `# contents of ./patterns/postfix:`
`POSTFIX_QUEUEID [0-9A-F]{10,11}`

grok

Jan 1 06:25:43 mailserver14 postfix/cleanup[21403]: BEF25A72965: message-id=<20130101142543.5828399CCAF@mailserver14.example.com>

```
filter {  
  grok {  
    patterns_dir => "./patterns"  
    match => { "message" => "%{SYSLOGBASE}  
%{POSTFIX_QUEUEID:queue_id}: %{GREEDYDATA:syslog_message}" }  
  }  
}
```

grok

- ▶ timestamp: Jan 1 06:25:43
- ▶ logsource: mailserver14
- ▶ program: postfix/cleanup
- ▶ pid: 21403
- ▶ queue_id: BEF25A72965
- ▶ syslog_message: message-id=<20130101142543.5828399CCAF@mailserver14.example.com>

date

Convert string-based date formats to date object for easy conversion and export.

date

- ▶ syslog events usually have timestamps like this:
`Apr 17 09:32:01`
- ▶ You would use the date format `MMM dd HH:mm:ss` to parse this.
- ▶ <http://www.joda.org/joda-time/apidocs/org/joda/time/format/DateTimeFormat.html>
- ▶ Overwrites `@timestamp` by default

date

```
filter {  
  # ...grok, etc.  
  date {  
    match => [ "timestamp", "MMM dd HH:mm:ss" ]  
    remove_field => { "timestamp" }  
    locale => "en"  
  }  
  # ...other filters  
}
```

date

- ▶ **ISO8601** - should parse any valid ISO8601 timestamp, such as 2011-04-19T03:44:01.103Z
- ▶ **UNIX** - will parse float or int value expressing unix time in seconds since epoch like 1326149001.132 as well as 1326149001
- ▶ **UNIX_MS** - will parse int value expressing unix time in milliseconds since epoch like 1366125117000
- ▶ **TAI64N** - will parse tai64n time values

geoip

Look up geographic information by
IP

geoip

```
geoip {  
  source => "clientip"  
}
```

useragent

Parse useragent strings into fields.

useragent

```
useragent {  
  source => "useragent"  
}
```

OR

```
if [useragent] != "" {  
  useragent { source => "useragent" }  
}
```

Filters

- aggregate
- alter
- anonymize
- collate
- csv
- cidr
- clone
- cipher
- checksum
- date
- dns
- drop
- elasticsearch
- extractnumbers
- environment
- elapsed
- fingerprint
- geoip
- grok
- i18n
- json
- json_encode
- kv
- mutate
- metrics
- multiline
- metaevent
- prune
- punct
- ruby
- range
- syslog_pri
- sleep
- split
- throttle
- translate
- uuid
- urldecode
- useragent
- xml
- zeromq

Conditionals

if/then/else

```
if EXPRESSION {  
    ...  
} else if EXPRESSION {  
    ...  
} else {  
    ...  
}
```

expressions

Comparison operators:

- equality: `==`, `!=`, `<`, `>`, `<=`, `>=`
- regexp: `=~`, `!~`
- inclusion: **`in`**, **`not in`**

Supported boolean operators:

- **`and`**, **`or`**, **`nand`**, **`xor`**

Supported unary operators:

- **`!`**

expressions

```
filter {  
  if [action] == "login" {  
    mutate { remove => "secret" }  
  }  
}
```

expressions

```
output {  
    # Send production errors to Zabbix  
    if [loglevel] == "ERROR" and [deployment] ==  
"production" {  
        zabbix {  
            ...  
        }  
    }  
}
```

expressions

```
if [foo] in [foobar] {  
  if [foo] in "foo" {  
    if "hello" in [greeting] {  
      if [foo] in ["hello", "world", "foo"] {  
        if [missing] in [alsomissing] {  
          if !("foo" in ["hello", "world"]) {
```

sprintf

- ▶ Reference field values within a string:

```
add_field => { "foo" => "%{bar}" }
```

```
add_field => { "foo_%{bar}" => "%{baz}" }
```

- ▶ Nested fields are referenced with square braces:

```
add_field => {  
  "foo" => "%{[@metadata][bar]}"  
}
```

zabbix

You know, for monitoring.

zabbix

- ▶ <https://github.com/logstash-plugins/logstash-output-zabbix>
- ▶ <https://www.elastic.co/guide/en/logstash/current/plugins-outputs-zabbix.html>
- ▶ Community plugin
- ▶ Deterministic (derives Zabbix host and key values from events)
- ▶ Installation:

```
bin/plugin install logstash-output-zabbix
```

zabbix

- ▶ `zabbix_sender` protocol
- ▶ Uses `@timestamp`
- ▶ Supports sending multiple values per event (most recently added feature)
- ▶ Uses native ruby TCP calls (old version used `zabbix_sender` binary)
- ▶ Does not support batching (don't overload your trappers)

options

- `zabbix_host`
- `zabbix_key`
- `zabbix_value`
- `zabbix_server_host`
- `zabbix_server_port`
- `multi_value`
- `timeout`

zabbix_host

- ▶ Type: String
- ▶ A single field name which holds the value you intend to use as the Zabbix host name.
- ▶ Required value.

zabbix_key

- ▶ Type: String
- ▶ A single field name which holds the value you intend to use as the Zabbix item key.
- ▶ Ignored if using `multi_value`, otherwise required.

zabbix_value

- ▶ Type: String
- ▶ A single field name which holds the value you intend to send to `zabbix_host`'s `zabbix_key`.
- ▶ Default: "message" (the whole, original log line)
- ▶ Ignored if using `multi_value`, otherwise required.

server

- `zabbix_server_host`

The IP or resolvable hostname where the Zabbix server is running

Default: `"localhost"`

- `zabbix_server_port`

The port on which the Zabbix server is running

Default: `10051`

multi_value

- ▶ Type: Array
- ▶ Ignores `zabbix_key` and `zabbix_value`.
- ▶ This can be visualized as:
 `[key1, value1, key2, value2, ... keyN, valueN]`
- ▶ ...where `key1` is an instance of `zabbix_key`, and `value1` is an instance of `zabbix_value`.
- ▶ If the field referenced by any `zabbix_key` or `zabbix_value` does not exist, that entry will be ignored.

timeout

- ▶ Type: Number
- ▶ The number of seconds to wait before giving up on a connection to the Zabbix server.
- ▶ Default: 1
- ▶ This number should be very small, otherwise delays in delivery of other outputs could result.

zabbix

```
output {  
    zabbix {  
        zabbix_server_host => "zabbix.example.com"  
        zabbix_host => "host_field"  
        zabbix_key => "key_field"  
        zabbix_value => "value_field"  
    }  
    # ... Other outputs  
}
```

zabbix

```
output {  
  if [type] == "zabbix" {  
    zabbix {  
      zabbix_server_host => "zabbix.example.com"  
      zabbix_host => "host_field"  
      zabbix_key => "key_field"  
      zabbix_value => "value_field"  
    }  
  }  
}
```

zabbix

```
output {  
  if [type] == "zabbix" {  
    zabbix {  
      zabbix_server_host => "zabbix.example.com"  
      zabbix_host => "host_field"  
      multi_value => [ "k1", "v1", "k2", "v2" ]  
    }  
  }  
}
```

use cases

It's play time!

IRC

- ▶ Monitor IRC for catch word or phrase
- ▶ Send to Zabbix if the word is given

input

```
input {  
  irc {  
    channels => [ "#zabbix" ]  
    host => "irc.freenode.org"  
    nick => "howdy"  
    port => 6667  
    type => "irc"  
  }  
}
```

filter

```
if [type] == "irc" {  
  if [message] =~ /^.*TESTING.*$/ {  
    mutate {  
      add_field => { "[@metadata][irc_key]" =>  
"message" }  
      add_field => { "[@metadata][zabbix_host]" =>  
"irc" }  
      add_tag => "testing"  
    }  
  }  
}
```


output

```
if [type] == "irc" and "testing" in [tags] {  
  zabbix {  
    zabbix_server_host => "localhost"  
    zabbix_host => "[@metadata][zabbix_host]"  
    zabbix_key => "[@metadata][irc_key]"  
    zabbix_value => "message"  
  }  
}
```

Result

Input (IRCCloud)

18:24:26 untergeek TESTING...one, two, three...

18:24:48 untergeek woot

18:25:03 untergeek |

Output (Zabbix Frontend)

<u>Last check</u>	Last value
2015-08-18 18:24:26	TESTING...one, two, three...

NGINX

- ▶ Capture NGINX logs for virtual hosts
- ▶ Watch for error codes (400 - 599)
- ▶ Send to Zabbix when one comes in
- ▶ Bonus: Send the client IP that generated the code

input

```
input {  
  file {  
    path => "/path/to/nginx.log"  
    type => "nginx_json"  
  }  
}
```

filter - pt.1

```
json {  
    source => "message"  
    remove_field => "message"  
}  
  
if [type] == "nginx_json" {  
    mutate {  
        replace => { "host" => "%{vhost}" }  
        remove_field => "vhost"  
    }  
}
```

filter - pt.2

```
geoip { source => "clientip" }  
if [useragent] != "" {  
    useragent { source => "useragent" }  
}  
if [referrer] == "-" {  
    mutate { remove_field => "referrer" }  
}
```

filter - pt.3

```
if [status] >= 400 and [host] != "localhost" {  
  mutate {  
    add_field => {  
      "[@metadata][status_key]" => "status"  
    }  
    add_field => {  
      "[@metadata][clientip_key]" => "clientip"  
    }  
  }  
}
```

filter - pt.4

```
add_field => {  
  "[@metadata][error]" => "error[%{status},]"  
}
```

```
add_field => {  
  "[@metadata][counter]" => "1"  
}
```

```
}
```

```
}
```

```
}
```


output - 1

```
if [type] == "nginx_json" {  
  if [status] >= 400 {  
    zabbix {  
      zabbix_server_host => "localhost"  
      zabbix_host => "host"  
      zabbix_key => "[@metadata][error]"  
      zabbix_value => "[@metadata][counter]"  
    }  
  }  
}
```



	zabbix host	key	value
fieldname	host	[@metadata][error]	[@metadata][counter]
value	<u>untergeek.com</u>	error[404,]	1

output - 2

```
zabbix {  
  zabbix_server_host => "localhost"  
  zabbix_host => "host"  
  multi_value => [  
    "[@metadata][status_key]", "status",  
    "[@metadata][clientip_key]", "clientip"  
  ]  
}
```

Result

- ▶ Two kinds here:

<u>Name</u> 	<u>Last check</u>	<u>Last value</u>
IP (1 Item)		
Client IP	2015-08-18 21:42:13	91.200 
Status (5 Items)		
HTTP Error Code 404	2015-08-18 21:16:08	1
HTTP Error Code 405	2015-08-18 19:53:50	1
HTTP Error Code 429	2015-08-18 20:40:00	1
HTTP Error Code 499	2015-08-18 21:42:13	1
HTTP Status	2015-08-18 21:42:13	499

Result

Timestamp	Value	Timestamp	Value
2015-08-18 21:42:13	91.200.	2015-08-18 21:42:13	499
2015-08-18 21:29:57	91.200.	2015-08-18 21:29:57	499
2015-08-18 21:26:46	91.200.	2015-08-18 21:26:46	499
2015-08-18 21:24:32	91.200.	2015-08-18 21:24:32	499
2015-08-18 21:21:38	52.10.2	2015-08-18 21:21:38	499
2015-08-18 21:16:08	66.249.	2015-08-18 21:16:08	404
2015-08-18 21:15:55	66.249.	2015-08-18 21:15:55	404
2015-08-18 21:14:25	66.249.	2015-08-18 21:14:25	404
2015-08-18 21:14:04	123.125	2015-08-18 21:14:04	404
2015-08-18 21:10:01	91.200.	2015-08-18 21:10:01	499

Result

► Just 404s

Timestamp	Value
2015-08-18 21:55:47	1
2015-08-18 21:16:08	1
2015-08-18 21:15:55	1
2015-08-18 21:14:25	1
2015-08-18 21:14:04	1

Conclusion

- ▶ <https://www.elastic.co/guide/en/logstash/current/index.html>
- ▶ <https://github.com/elastic/logstash>
- ▶ <https://github.com/logstash-plugins/logstash-output-zabbix>
- ▶ <https://discuss.elastic.co/c/logstash>
- ▶ #logstash on [irc.freenode.org](https://www.freenode.net/)